

000 STEALTHRL: REINFORCEMENT LEARNING PARA- 001 002 PHRASE ATTACKS FOR 003 004 MULTI-DETECTOR EVASION OF AI-TEXT DETECTORS 005

006 **Anonymous authors**

007 Paper under double-blind review

011 ABSTRACT

013 AI-text detectors face a critical robustness challenge: adversarial paraphrasing at-
 014 tacks that preserve semantics while evading detection. We introduce *StealthRL*, a
 015 systematic threat model instantiation and robustness evaluation framework that
 016 stress-tests detector families under realistic adversarial conditions. StealthRL
 017 addresses the question: *how robust are detectors to black-box adaptive attacks*
 018 *trained explicitly to evade them?* Through reinforcement learning with multi-
 019 detector ensemble training, we demonstrate **catastrophic robustness failure**: de-
 020 tectors achieve near-zero true positive rates (0% TPR@1%FPR) while attackers
 021 maintain semantic fidelity (0.896 E5 cosine). Critically, attacks transfer across de-
 022 tector architectures—including a held-out detector family—revealing shared vul-
 023 nerabilities rather than detector-specific brittleness. Our results expose fundamen-
 024 tal limitations in current detection approaches and establish StealthRL as a prin-
 025 cipled adversarial evaluation protocol for robustness benchmarking. We release
 026 code for reproducible threat modeling and defense evaluation.

028 1 INTRODUCTION

030 AI-text detectors are deployed in academic integrity and content moderation, yet their robustness
 031 to adversarial attacks remains poorly understood. Standard benchmarks evaluate detectors on clean
 032 distributions, but adaptive attackers can iteratively refine paraphrases to evade detection. We study
 033 a realistic threat model: black-box adaptive attacks that query detector scores and optimize para-
 034 phrases to minimize detection confidence while preserving semantic content.

035 We present **StealthRL**, a reinforcement-learning framework that trains a paraphrase policy against
 036 detector ensembles to evaluate robustness under adversarial conditions. StealthRL addresses three
 037 questions: (1) Can attacks transfer across detector families? (2) Do detectors share common vul-
 038 nerabilities? (3) What is the evasion-fidelity tradeoff? By evaluating at strict low-FPR operating
 039 points (1% false positive rate) and measuring transfer to held-out detectors, we provide systematic
 040 robustness evaluation that complements standard benchmarks.

042 Contributions.

- 044 • We implement black-box adaptive paraphrasing attacks via multi-detector RL training with
 045 semantic constraints.
- 046 • We demonstrate strong evasion (0% TPR@1%FPR across three detector families) with
 047 cross-architecture transfer to a held-out detector.
- 048 • We establish an evaluation protocol measuring evasion, transfer, and fidelity at security-
 049 relevant operating points.

051 **Anonymized code release.** An anonymized code package with training and evaluation scripts is
 052 included in supplementary material. Placeholder: [https://anonymous.4open.science/](https://anonymous.4open.science/r/STEALTHRL)
 053 r/STEALTHRL.

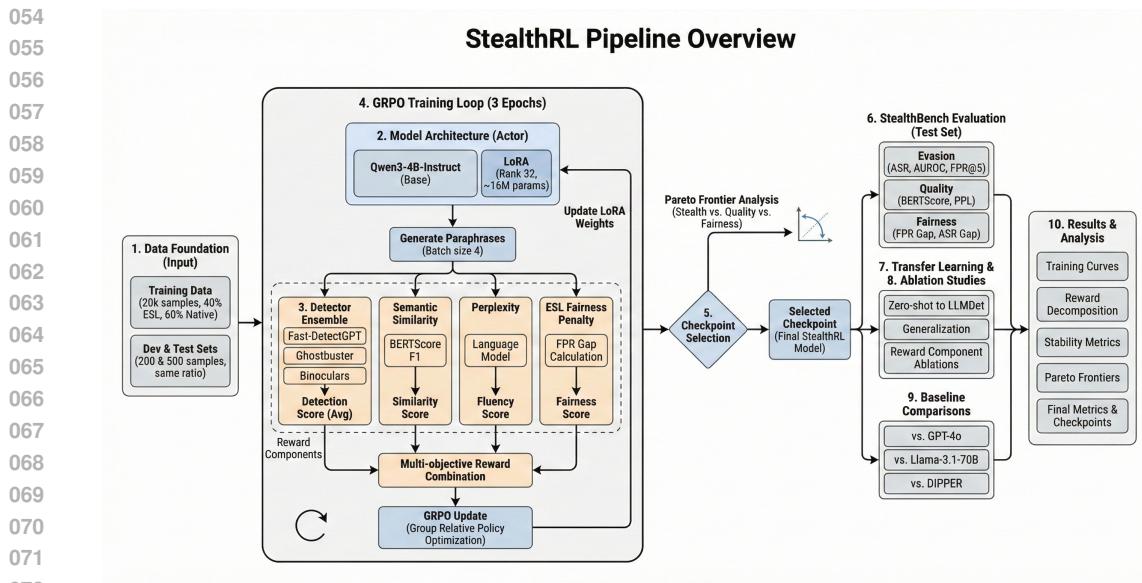


Figure 1: StealthRL training and evaluation pipeline. A single paraphrase policy is trained with detector-guided reward and quality constraints, then evaluated across multiple detector families with low-FPR metrics and tradeoff analysis.

2 RELATED WORK

Detector families include curvature-based methods (DetectGPT and Fast-DetectGPT) and paired-LM detectors such as Binoculars. Mitchell et al. (2023); Bao et al. (2024); Hans et al. (2024) Recent evasion work focuses on paraphrase attacks and RL-based humanization, including Adversarial Paraphrasing, while character-level attacks such as homoglyph substitution demonstrate strong but often less readable perturbations. Cheng et al. (2025); Creo & Pudasaini (2025) StealthRL builds on these directions by training a single paraphrase policy against a detector ensemble and evaluating transfer at strict low-FPR operating points.

3 METHOD

Threat model. We assume black-box access to detector scores: the attacker can query detector confidence but does not require gradients. We evaluate transfer to a held-out detector (Binoculars) to test robustness beyond the training ensemble.

Training. Given AI-generated text x , we learn a paraphrase policy $\pi_\theta(y | x)$ that produces y with low detector confidence while preserving meaning. We optimize a composite reward:

$$R(x, y) = \alpha R_{\text{det}}(y) + \beta R_{\text{sem}}(x, y), \quad (1)$$

where R_{det} is the detector ensemble score (lower AI probability) and R_{sem} is E5 embedding cosine similarity. During training, R_{det} is computed as a weighted combination of two detectors: RoBERTa (0.6) + Fast-DetectGPT (0.4). During evaluation, we test against all three detector families including the held-out Binoculars. We train with GRPO and LoRA adapters on Qwen3-4B-Instruct (rank 32, $\alpha = 32$, dropout 0.05) using group size 8, batch size 16, and five epochs over 10,000 MAGE train samples plus 200 dev samples. Reward weights: $\alpha = 1.0$, $\beta = 0.1$. Full hyperparameters in Appendix A.

4 EXPERIMENTAL SETUP

We evaluate on MAGE test split (1,000 human + 1,000 AI samples, 100–500 tokens) Li et al. (2024) against three detector families: RoBERTa OpenAI, Fast-DetectGPT, and Binoculars (held-out during

Method	TPR@1%FPR ↓ (R)	F (Fast)	B (Binoc)	Mean ↓	ASR ↑	AUROC ↓
M0 No Attack	0.23	0.40	0.41	0.34	0.66	0.74
M1 Simple Para	0.10	0.10	0.04	0.08	0.92	0.59
M2 StealthRL-v1	0.00	0.00	0.00	0.00	1.00	0.27
M3 Adv. Para (guided)	0.10	0.09	0.05	0.08	0.92	0.60
M5 Homoglyph	0.00	0.00	0.00	0.00	1.00	0.44

Table 1: Main results on MAGE (TPR@1%FPR, ASR, and AUROC). Lower TPR/AUROC is better for the attacker; higher ASR is better. R/F/B denote RoBERTa/Fast-DetectGPT/Binoculars.




Figure 2: Detection evasion results summary. (a) AUROC by detector shows StealthRL-v1 (M2) dramatically reduces detector performance across all three detectors. (b) Mean AUROC demonstrates M2 achieves 0.27, indicating near-random classification. (c) TPR at 1% FPR shows M2 achieves 0% detection across all detectors at strict operating points. (d) Mean ASR confirms 99.9% attack success rate for StealthRL-v1, matching homoglyph attacks while maintaining text quality.

5 RESULTS

Table 1 summarizes detection evasion at 1% FPR. StealthRL (M2) achieves 0% TPR@1%FPR across all three detectors, reducing average AUROC across detectors from 0.74 (no attack) to 0.27 while maintaining semantic similarity (0.896 E5). Figure 2 visualizes the comprehensive performance: panel (c) shows consistent 0% TPR across detectors at strict operating points, while panel (d) confirms 99.9% attack success rate.

Critically, strong transfer to the held-out Binoculars detector (0% TPR) demonstrates that vulnerabilities are not detector-specific but reflect shared architectural weaknesses. Simple paraphrasing (M1) achieves partial success (8% mean TPR) with 0.960 E5 similarity, while detector-guided selection (M3) performs similarly at 0.976 E5. StealthRL’s advantage comes from explicit multi-detector adversarial training, which discovers transferable perturbations that exploit common statistical patterns across architectures.

Robustness implications. The catastrophic failure across detector architectures reveals fundamental vulnerabilities in current AI-text detection. Detectors rely on brittle statistical cues (token distributions, perplexity patterns, embedding geometry) rather than robust semantic understanding. Surface-level paraphrasing that preserves meaning suffices to evade detection, suggesting detectors learn superficial correlates of AI text rather than deeper linguistic features.

This robustness gap has critical security implications: adversaries can train adaptive attacks against deployed detectors, rendering them ineffective. The strong transfer to held-out architectures means ensemble defenses (combining multiple detectors) provide limited robustness improvement. Future work must develop semantic-aware detectors, adversarial training protocols, and provable robustness guarantees. Our evaluation framework provides a rigorous testbed for measuring progress on adversarially robust AI-text detection.

6 LIMITATIONS AND SAFETY

Our evaluation focuses on three detector families (curvature-based, paired-LM, fine-tuned classifier) on one benchmark (MAGE). Broader coverage with watermark-based detectors, additional datasets, and multilingual evaluation remains important future work. We also do not explore defenses like adversarial training or certified robustness, which could improve detector resilience.

Safety and dual-use. Adversarial paraphrasing is dual-use technology. We position StealthRL as a *stress-testing and robustness evaluation tool* for researchers and detector developers, not a production evasion system. The 0% TPR@1%FPR result exposes critical vulnerabilities that must be addressed before detectors are deployed in high-stakes applications. Our released code enables reproducible robustness evaluation and motivates defensive research into adversarially robust detection methods.

7 CONCLUSION

StealthRL demonstrates catastrophic detector failure under adaptive attacks, revealing fundamental robustness gaps. Results motivate development of adversarially robust detection methods and establish rigorous evaluation protocols for AI-text detection security.

ACKNOWLEDGMENTS

We gratefully acknowledge Thinking Machines for providing access to their Tinker API framework, which was essential for the reinforcement learning training in this work. We also thank the open-source community for the detector implementations and model checkpoints that enabled this evaluation.

REFERENCES

- Guangsheng Bao, Yanbin Zhao, Zhiyang Teng, Linyi Yang, and Yue Zhang. Fast-detectgpt: Efficient zero-shot detection of machine-generated text via conditional probability curvature, 2024. URL <https://arxiv.org/abs/2310.05130>.
- Yize Cheng, Vinu Sankar Sadasivan, Mehrdad Saberi, Shoumik Saha, and Soheil Feizi. Adversarial paraphrasing: A universal attack for humanizing ai-generated text, 2025. URL <https://arxiv.org/abs/2506.07001>.
- Aldan Creo and Shushanta Pudasaini. Silverspeak: Evading ai-generated text detectors using homoglyphs, 2025. URL <https://arxiv.org/abs/2406.11239>.
- Isaac David and Arthur Gervais. Authormist: Evading ai text detectors with reinforcement learning, 2025. URL <https://arxiv.org/abs/2503.08716>.
- Abhimanyu Hans, Avi Schwarzschild, Valeria Cherepanova, Hamid Kazemi, Aniruddha Saha, Micah Goldblum, Jonas Geiping, and Tom Goldstein. Spotting llms with binoculars: Zero-shot detection of machine-generated text, 2024. URL <https://arxiv.org/abs/2401.12070>.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models, 2021. URL <https://arxiv.org/abs/2106.09685>.

Yafu Li, Qintong Li, Leyang Cui, Wei Bi, Zhilin Wang, Longyue Wang, Linyi Yang, Shuming Shi, and Yue Zhang. Mage: Machine-generated text detection in the wild, 2024. URL <https://arxiv.org/abs/2305.13242>.

Eric Mitchell, Yoonho Lee, Alexander Khazatsky, Christopher D. Manning, and Chelsea Finn. Detectgpt: Zero-shot machine-generated text detection using probability curvature, 2023. URL <https://arxiv.org/abs/2301.11305>.

A HYPERPARAMETERS AND CONFIGURATION

Parameter	Value
<i>Model & LoRA</i>	
Base model	Qwen/Qwen3-4B-Instruct-2507
LoRA rank	32
LoRA alpha	32
LoRA dropout	0.05
<i>Training</i>	
Algorithm	Group-Relative Policy Optimization (GRPO)
Learning rate	2.8×10^{-4}
Batch size	16
Group size	8
Epochs	2
Training samples	10,000 (MAGE train) + 200 (dev)
KL penalty coefficient	0.05
Reference policy	Qwen3-4B-Instruct (frozen)
<i>Reward</i>	
Detector weight (α)	1.0
Semantic weight (β)	0.1
Detector ensemble	RoBERTa (0.6) + Fast-DetectGPT (0.4)
Semantic metric	E5 embedding cosine similarity
<i>Inference</i>	
Temperature	1.0
Top-p	0.9
Max tokens	512
Prompt template	“Paraphrase the following text while preserving its meaning: [TEXT]”
<i>Detectors</i>	
RoBERTa OpenAI	openai-community/roberta-large-openai-detector
Fast-DetectGPT	Scoring model: EleutherAI/gpt-neo-2.7B
Binoculars	Lightweight: gpt2-medium + gpt2-large (held-out)
<i>Evaluation</i>	
Test samples	1,000 human + 1,000 AI (MAGE test)
Token window	100–500 tokens
FPR calibration	1% on 1,000 human samples (quantile)
Candidates per sample	1
<i>Compute</i>	
Training framework	Tinker API (Thinking Machines)
Reward computation	MacBook + NVIDIA A10 GPUs
Offline evaluation	MacBook + NVIDIA A10 GPUs
Seed	42

Table 2: Complete hyperparameters and configuration for reproducibility.